# CS231n Lecture 5


- BOAZ 10기 박성현
- BOAZ 11기 김태희
- BOAZ 11기 홍지민

# Fully Connected Layer

32x32x3 image -> stretch to 3072 x 1

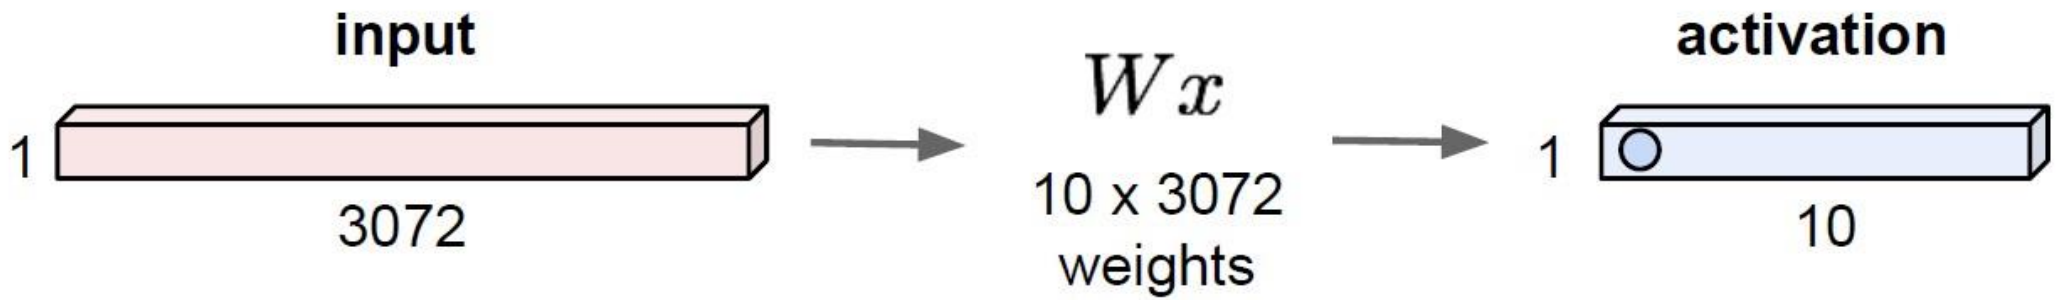

# Convolution Layer

32x32x3 image -> preserve spatial structure

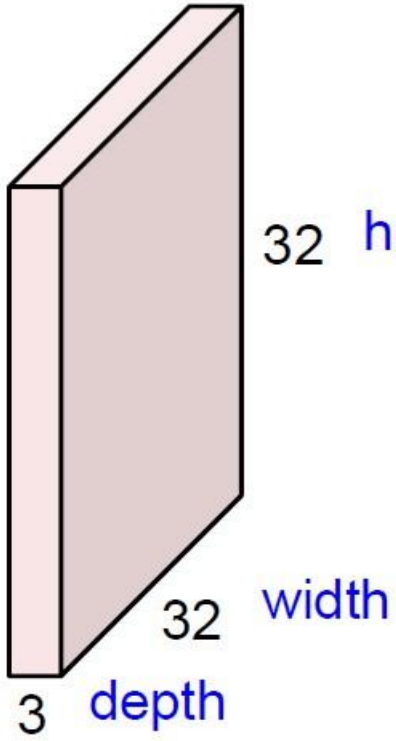

# Dot product & filter size

32x32x3 image

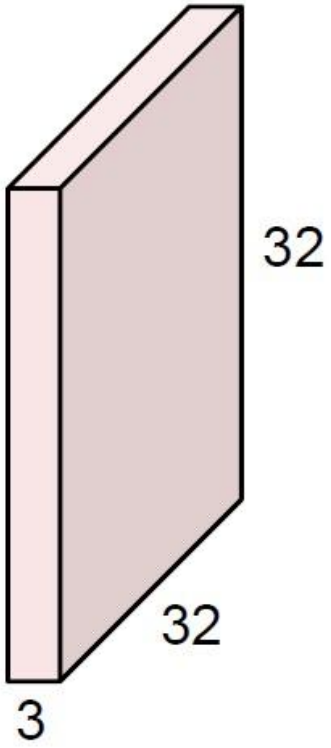


5x5x3 filter

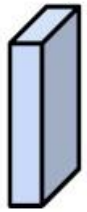

**Convolve** the filter with the image
i.e. “slide over the image spatially, computing dot products”

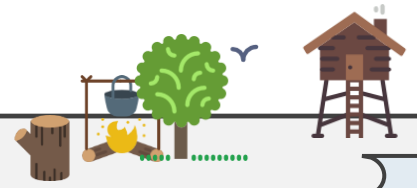

# Convolution Layer

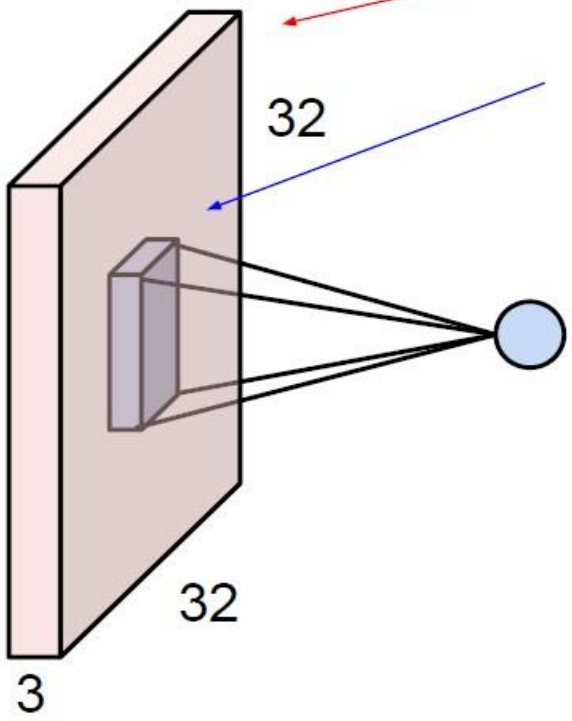


**1 number:**
the result of taking a dot product between the filter and a small 5x5x3 chunk of the image
(i.e. 5*5*3 = 75-dimensional dot product + bias)

$$w^T x + b$$

# Convolution Layer

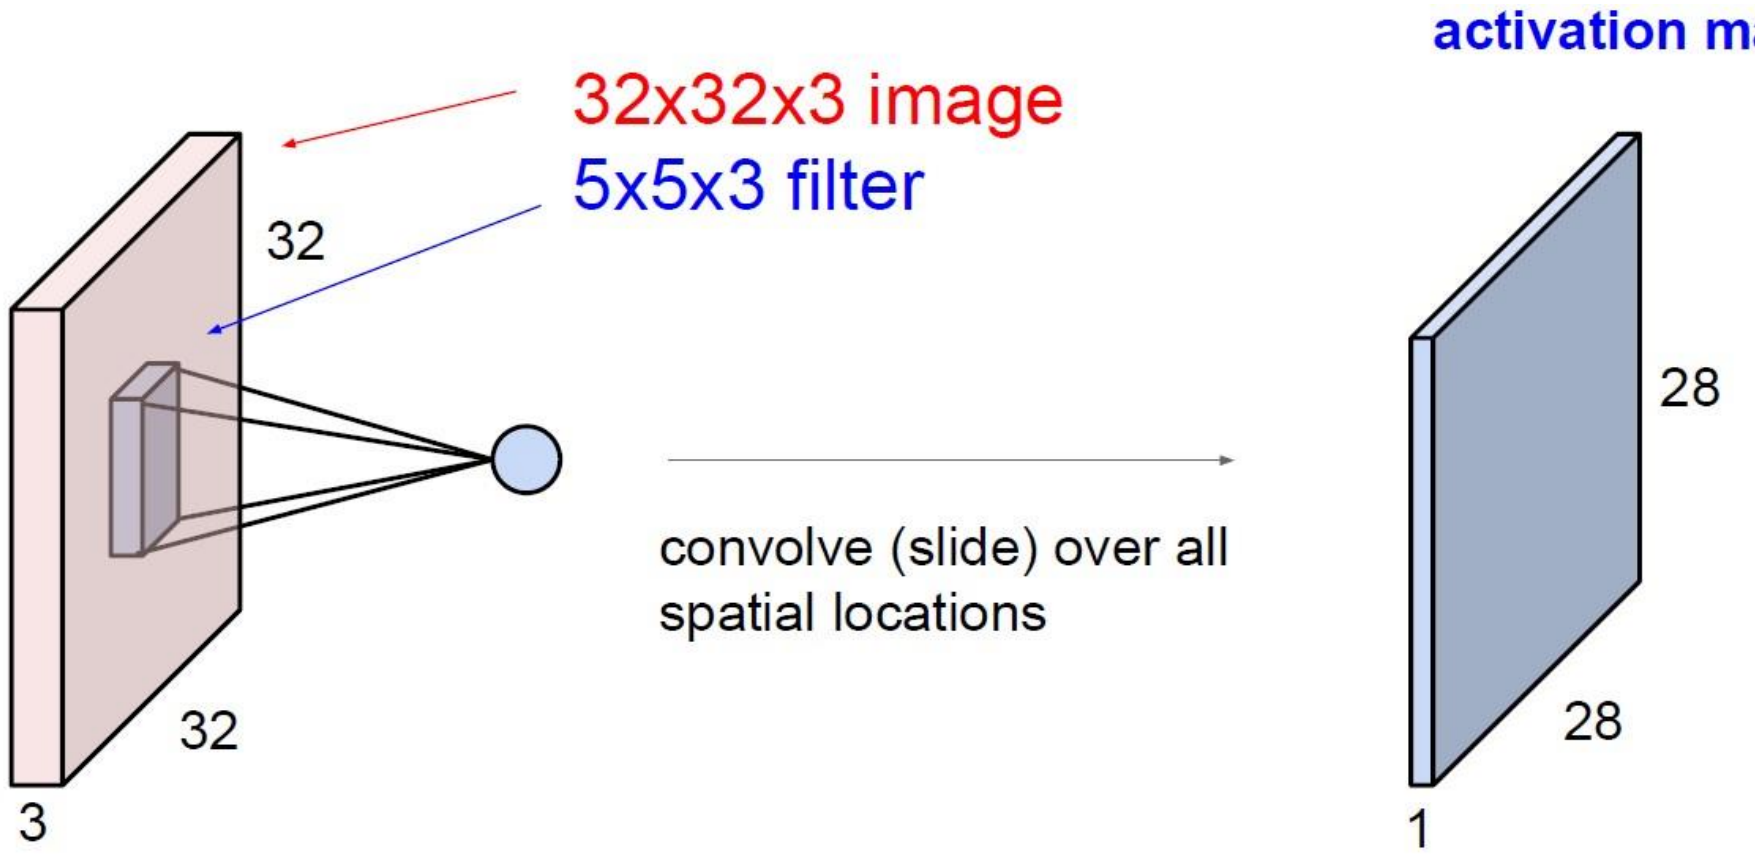

# Convolution Layer

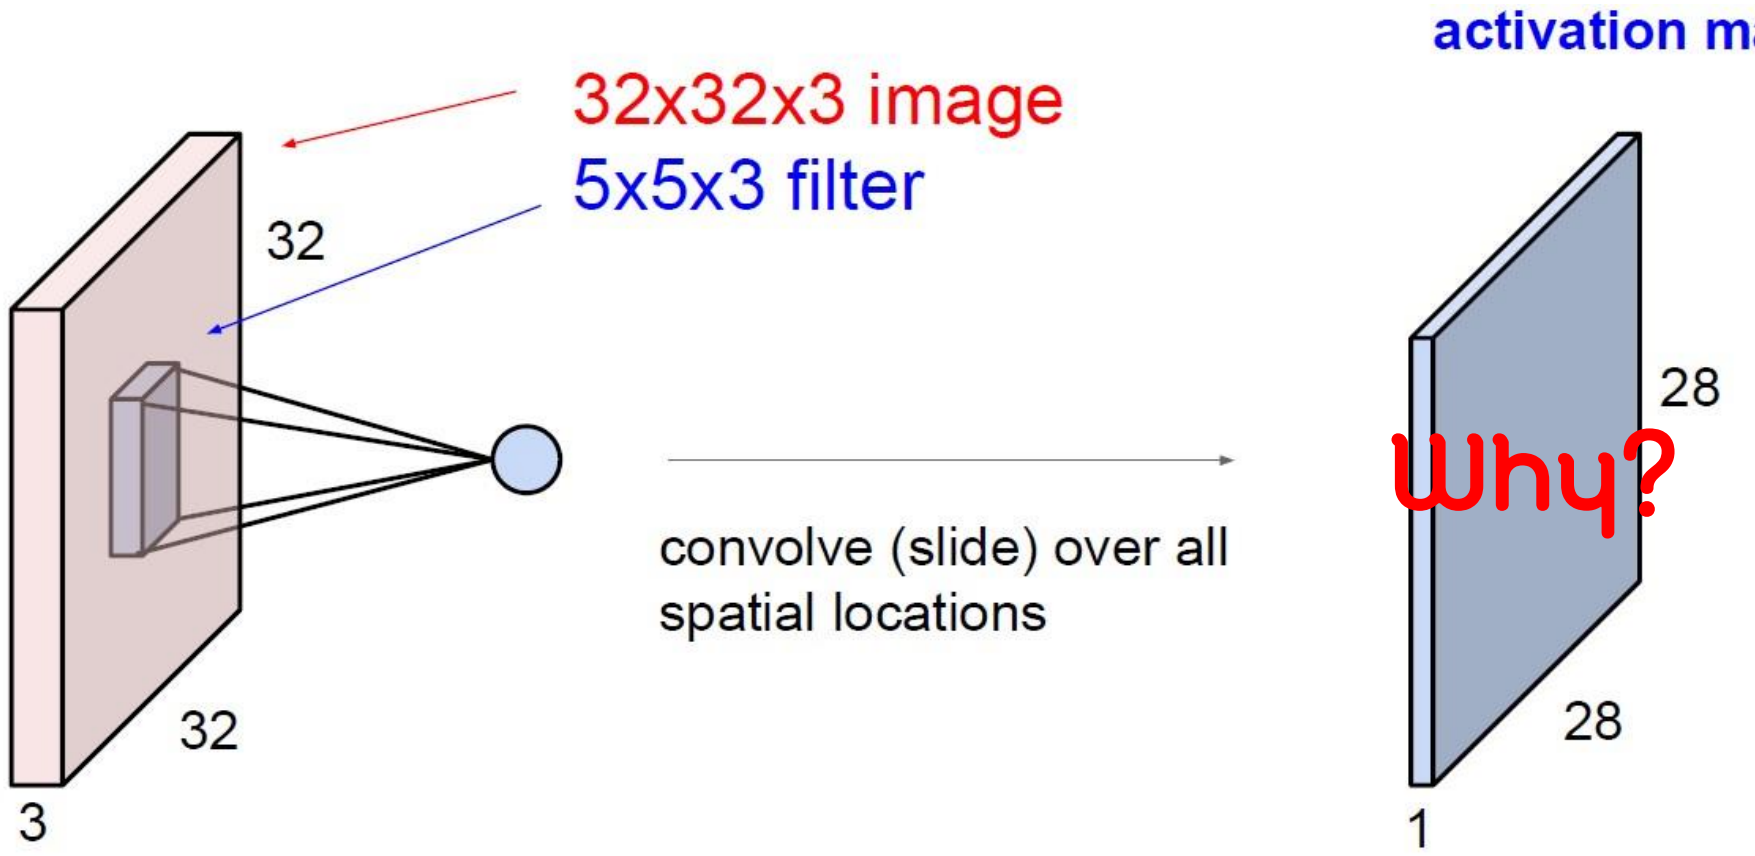

# Filter computation

A closer look at spatial dimensions:

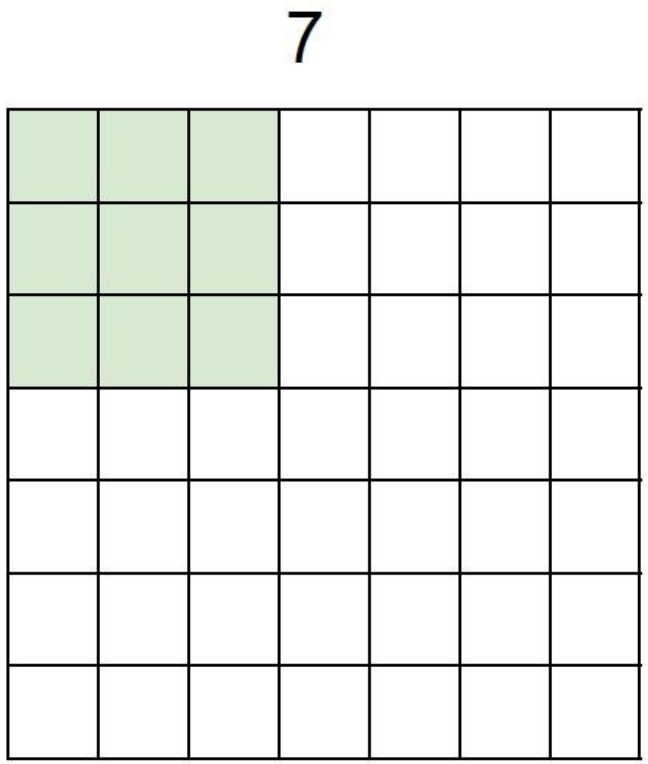


7x7 input (spatially)
assume 3x3 filter

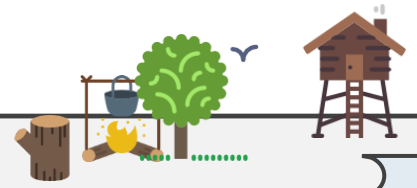

# Filter computation

A closer look at spatial dimensions:

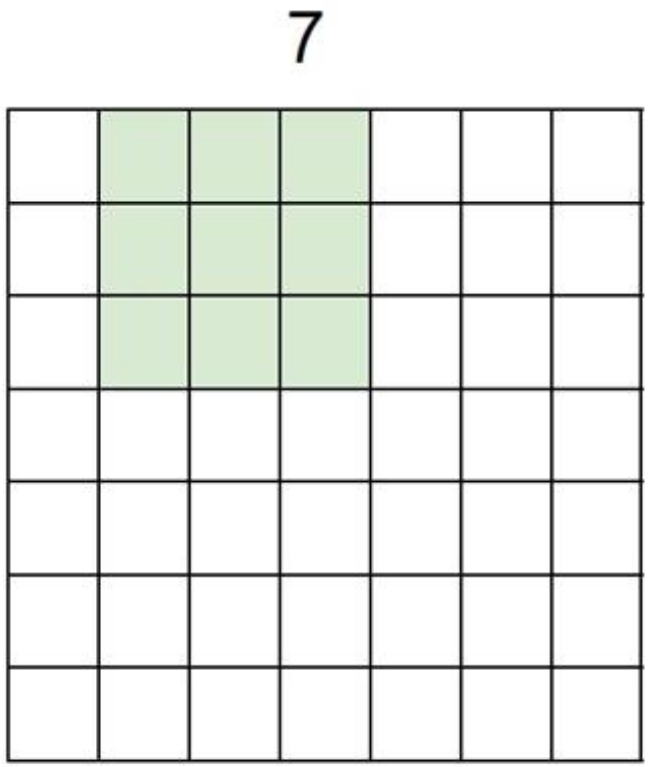


7x7 input (spatially)
assume 3x3 filter

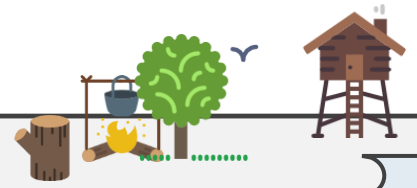

# Filter computation

A closer look at spatial dimensions:

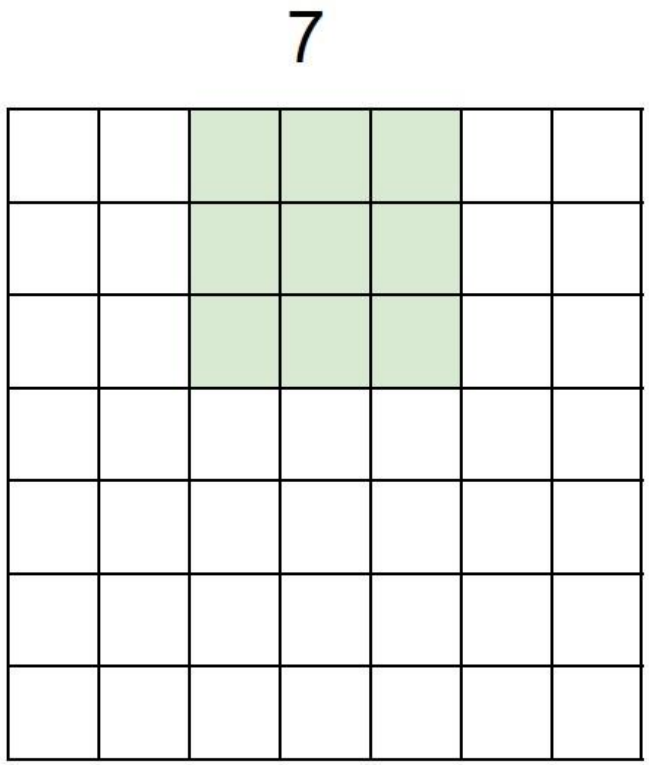


7x7 input (spatially)
assume 3x3 filter

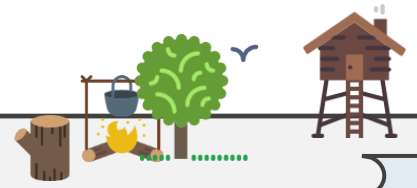

# Filter computation

A closer look at spatial dimensions:

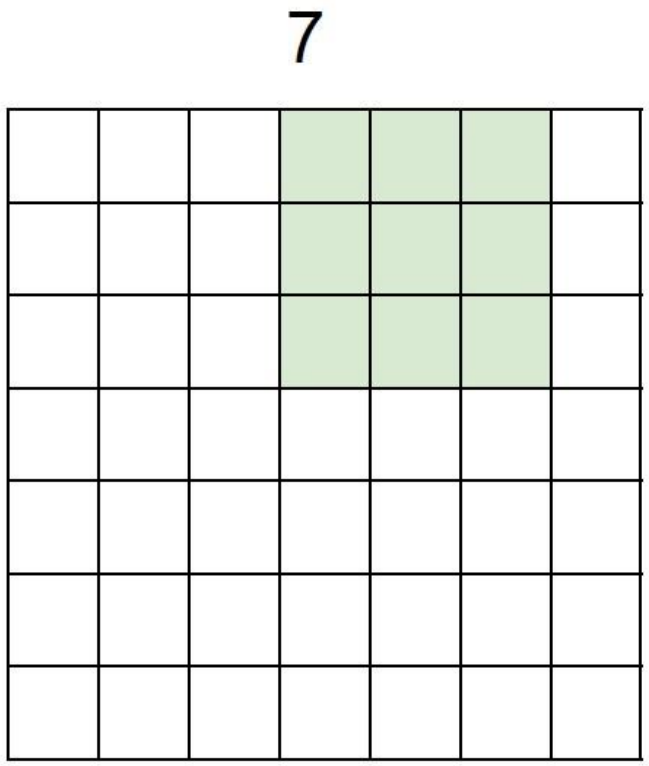


7x7 input (spatially)
assume 3x3 filter

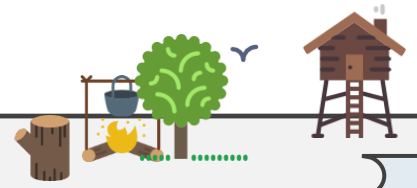

# Filter computation

A closer look at spatial dimensions:

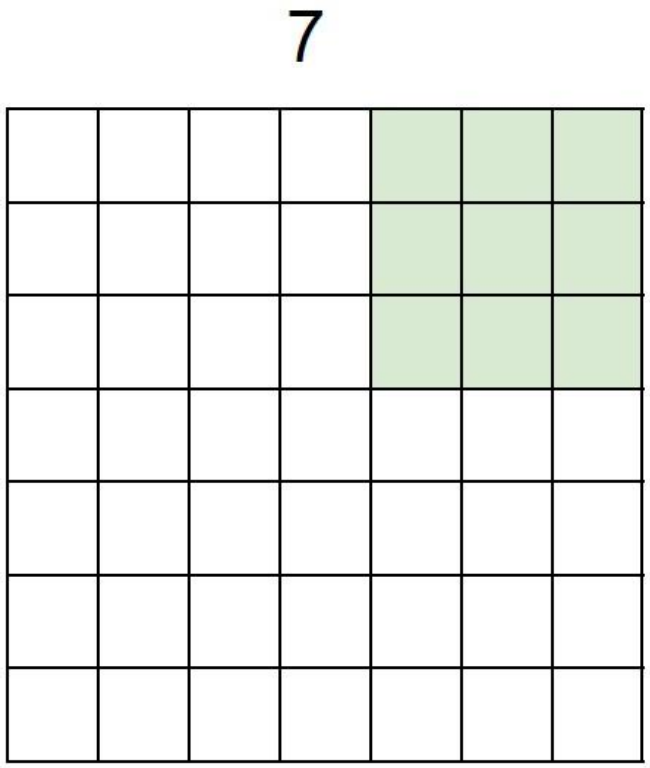


7x7 input (spatially)
assume 3x3 filter

**=> 5x5 output**

# Filter computation

A closer look at spatial dimensions:

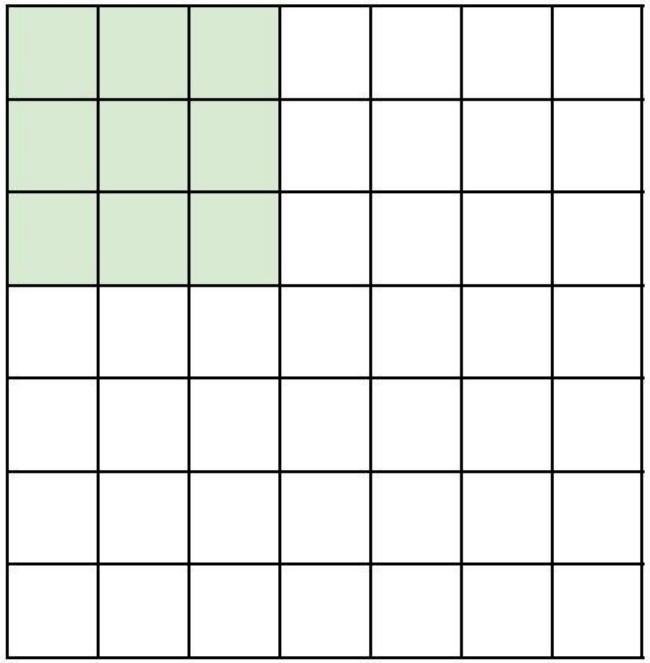


7x7 input (spatially)
assume 3x3 filter
applied **with stride 2**

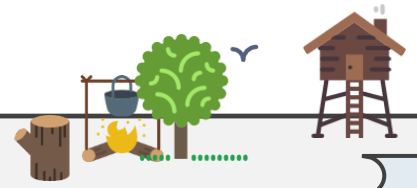

# Filter computation

A closer look at spatial dimensions:

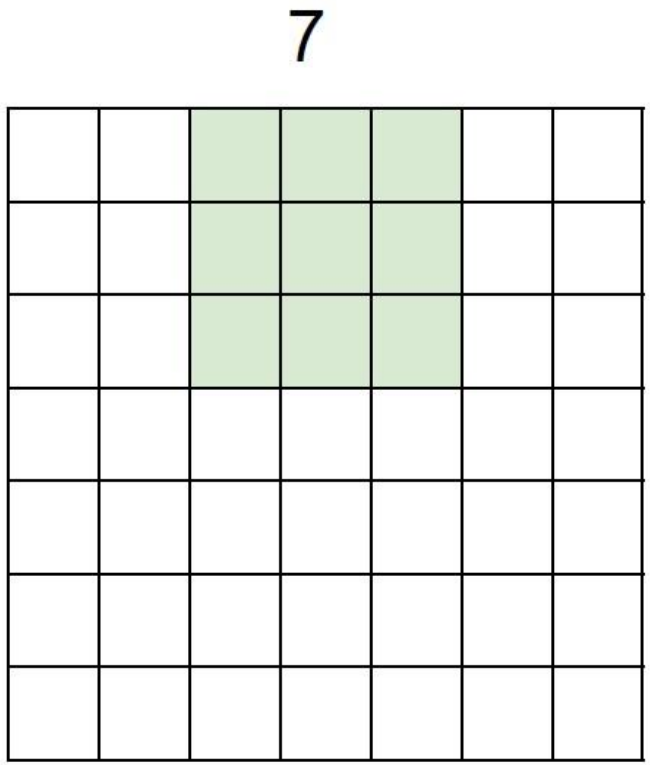


7x7 input (spatially)
assume 3x3 filter
applied **with stride 2**

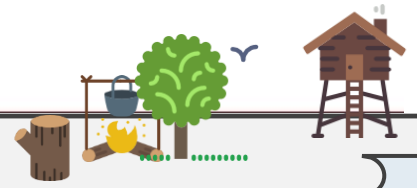

# Filter computation

A closer look at spatial dimensions:

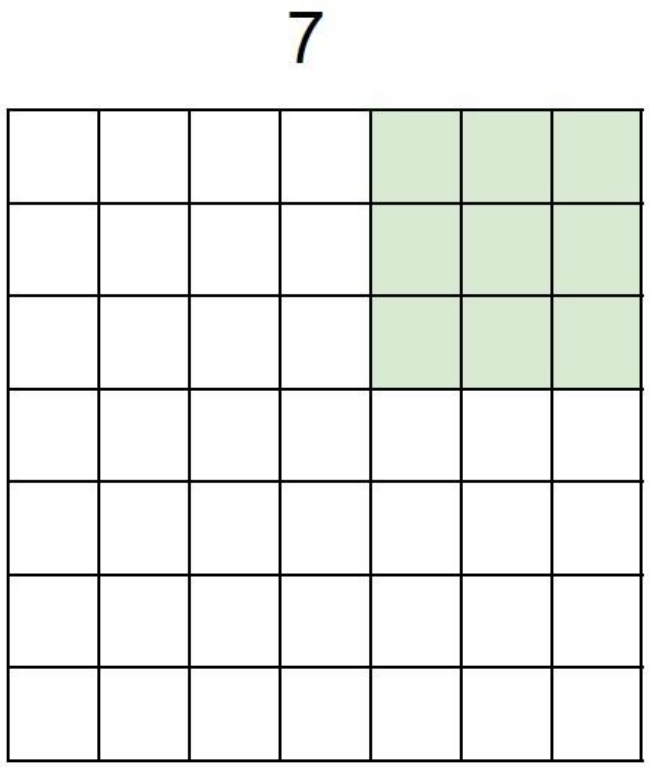


7x7 input (spatially)
assume 3x3 filter
applied **with stride 2**
**=> 3x3 output!**

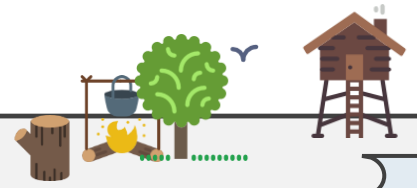

# Filter computation

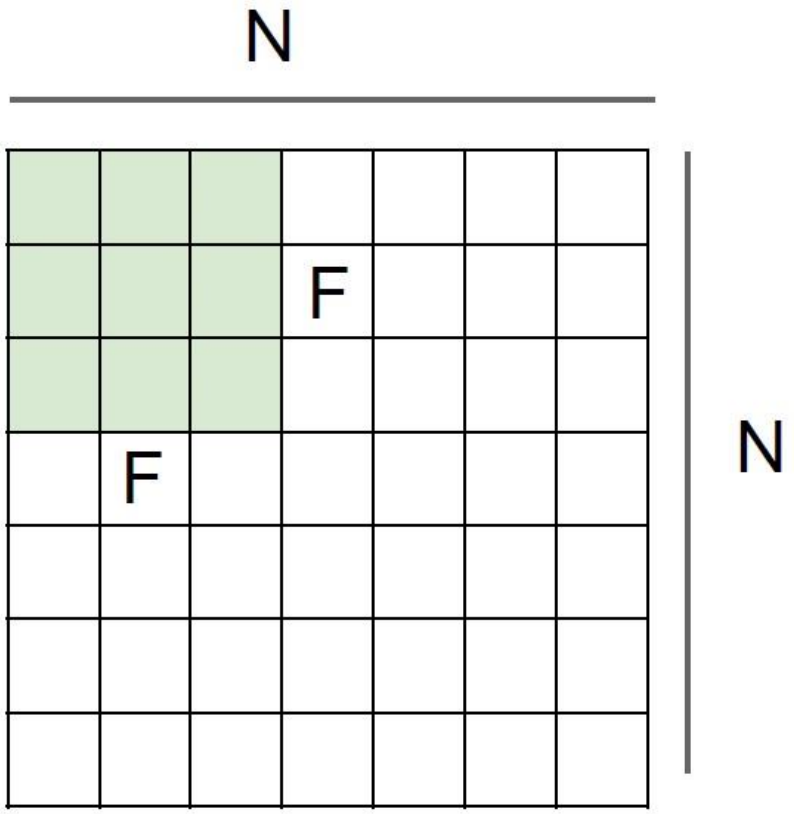


Output size:

**(N - F) / stride + 1**

e.g. N = 7, F = 3:

stride 1 => (7 - 3)/1 + 1 = 5

stride 2 => (7 - 3)/2 + 1 = 3

stride 3 => (7 - 3)/3 + 1 = 2.33 :\

# Convolution Layer computation

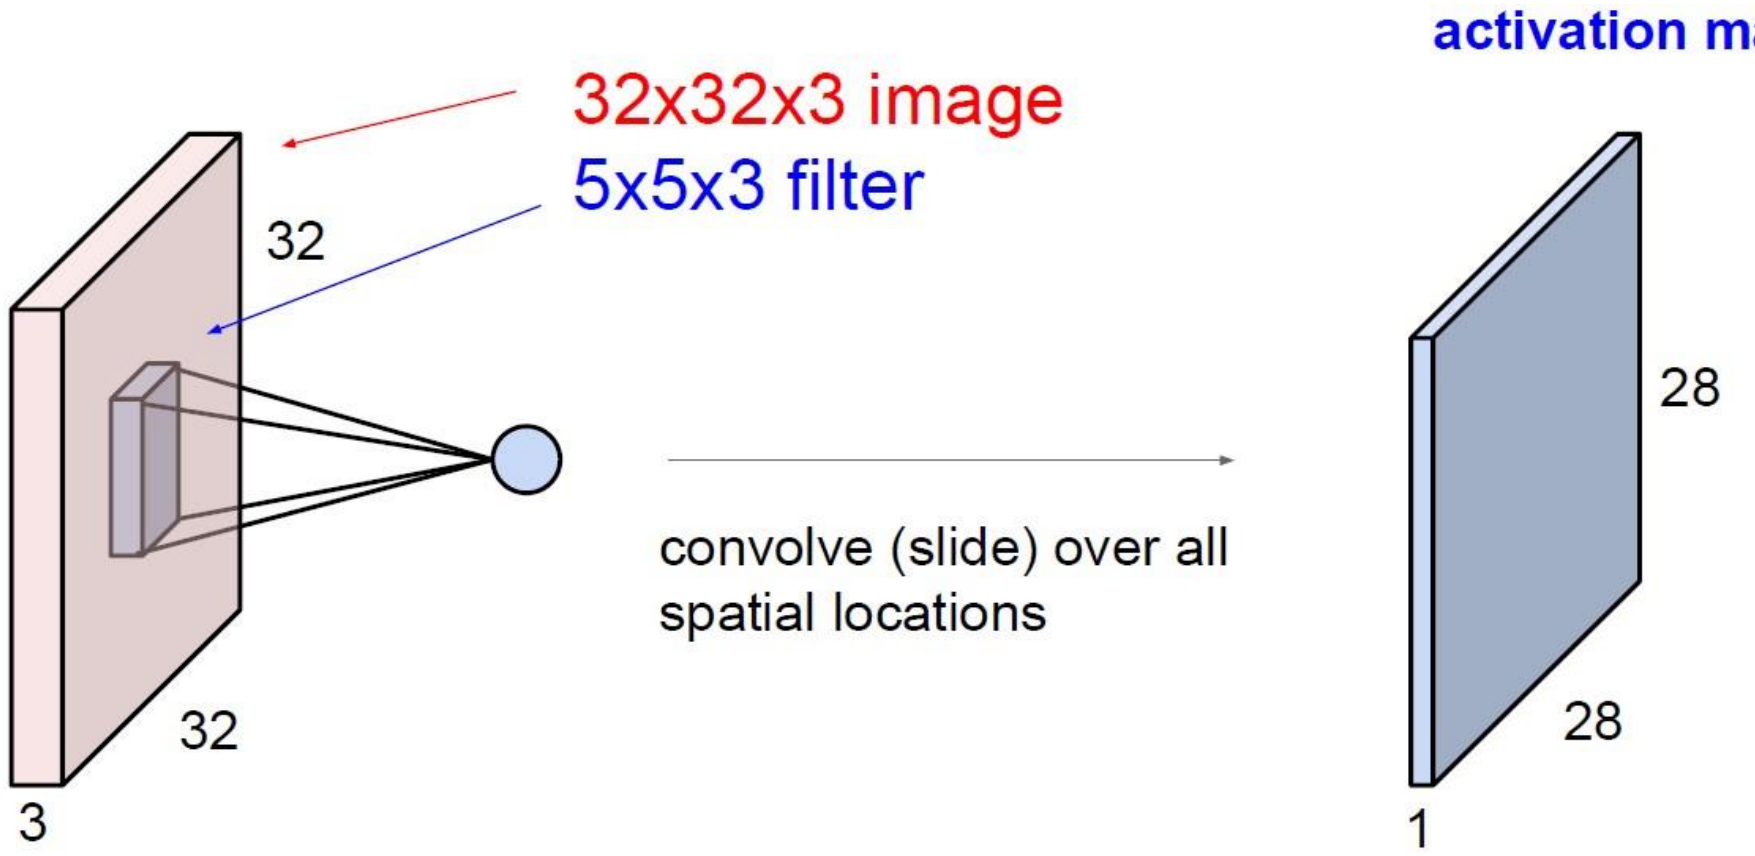


**(N - F) / stride + 1** (32−5)/1+1 = 28

# Convolution Layer computation

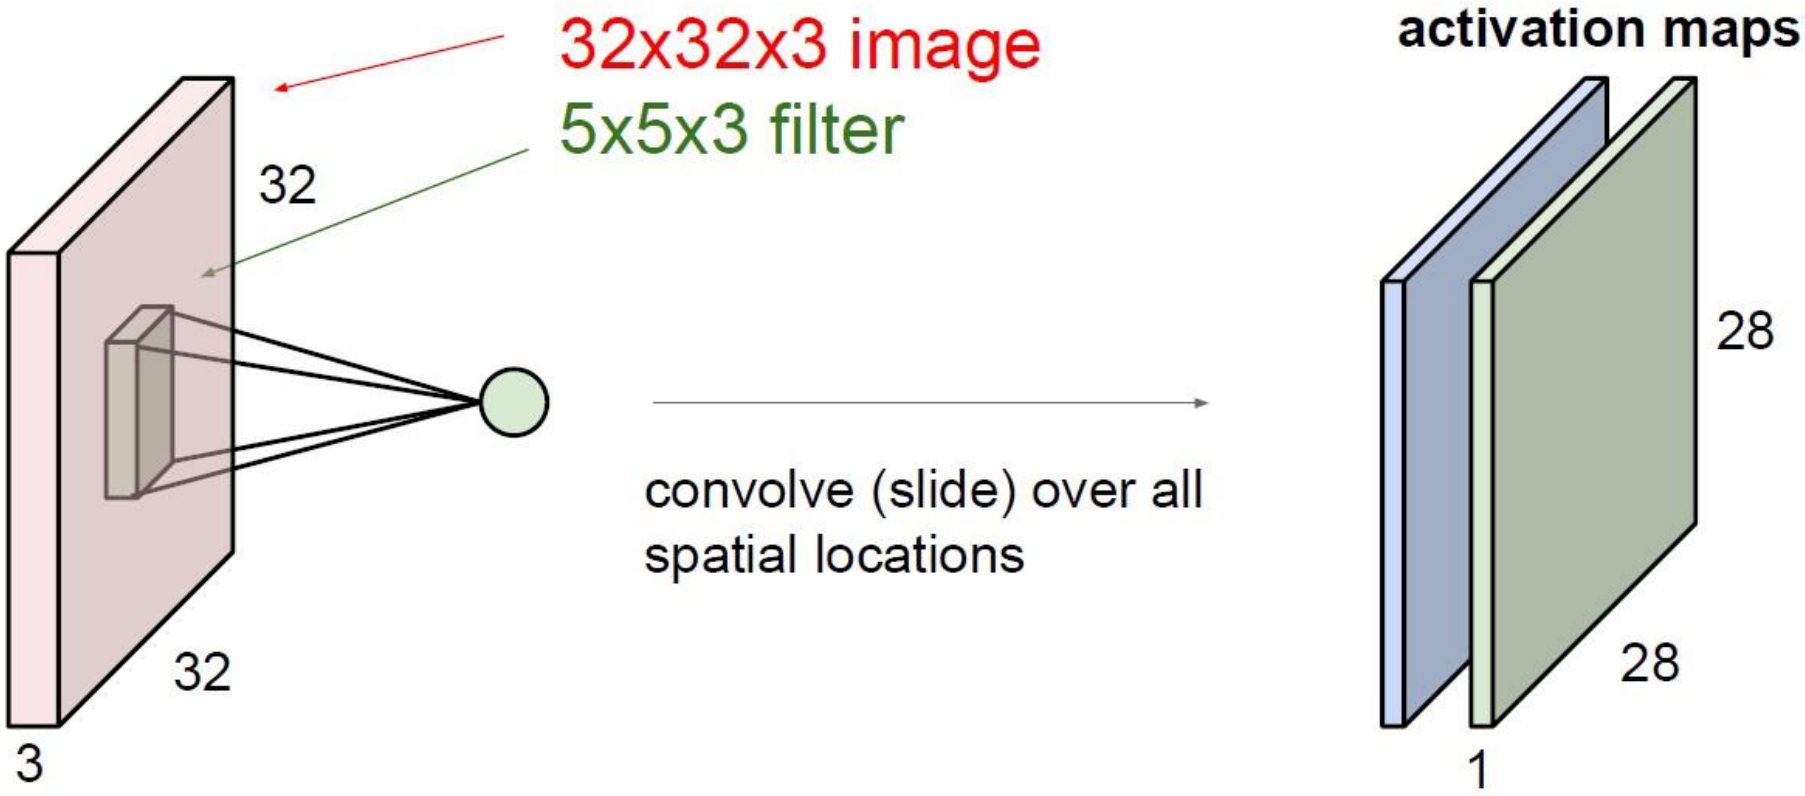

# Convolution Layer computation

For example, if we had 6 5x5 filters, we'll get 6 separate activation maps:

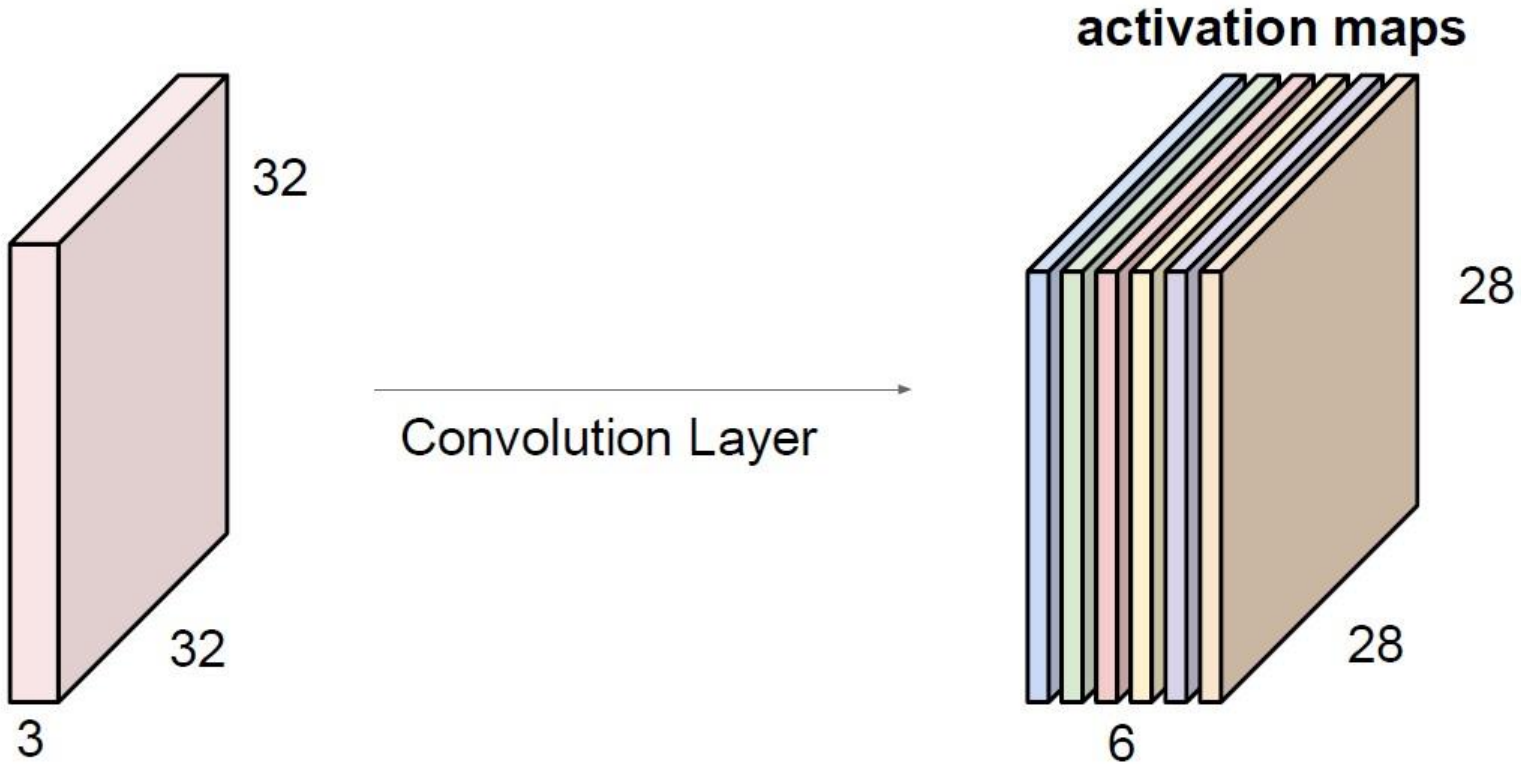


We stack these up to get a "new image" of size 28x28x6!

# Convolution Network

**Preview:** ConvNet is a sequence of Convolution Layers, interspersed with activation functions

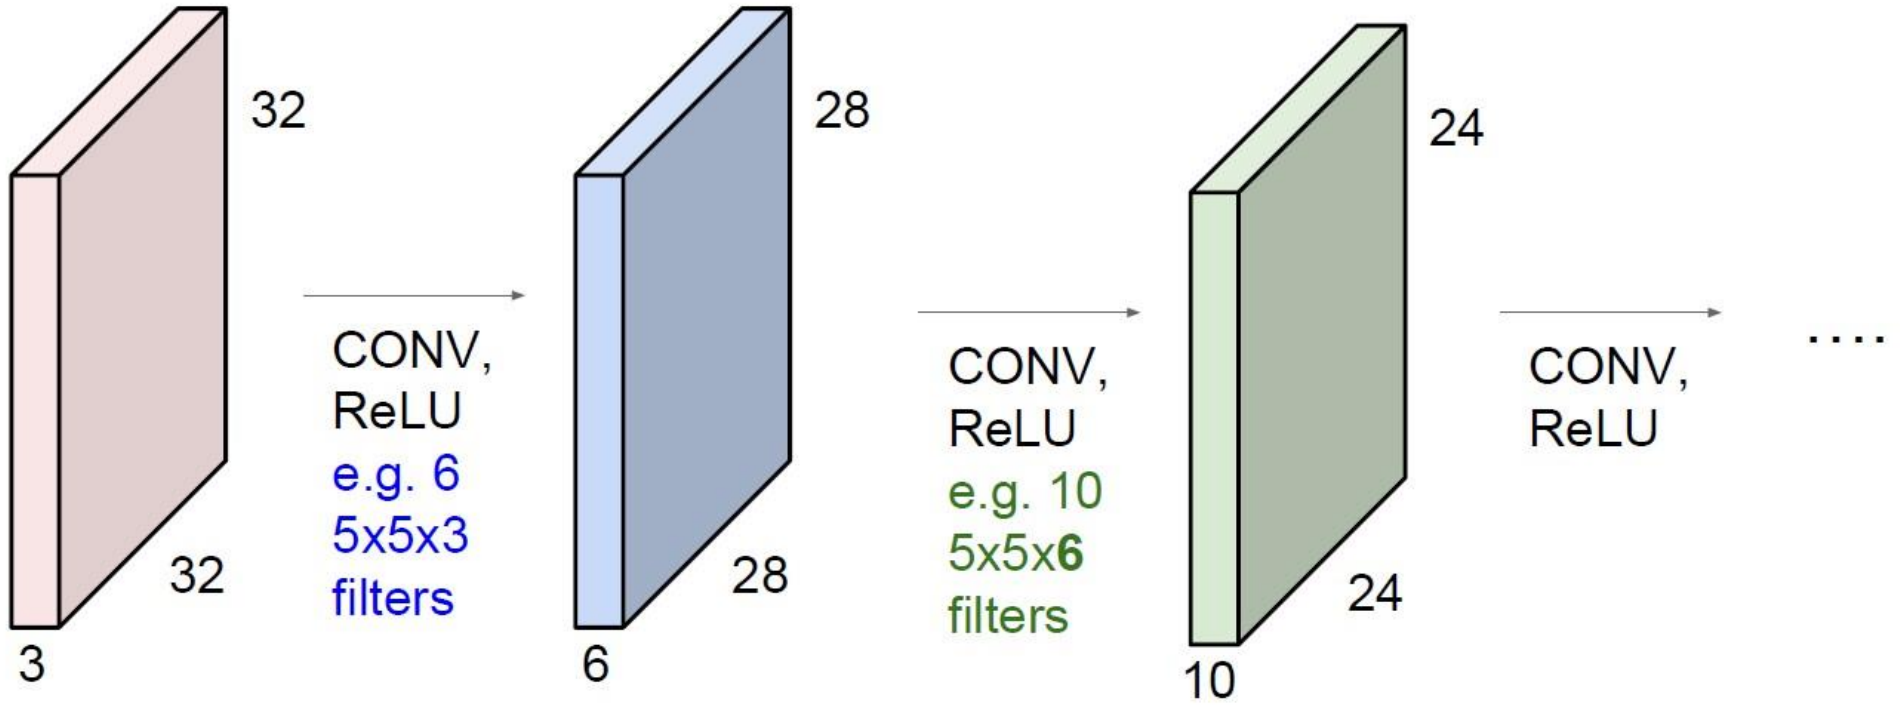


# Convolution Network

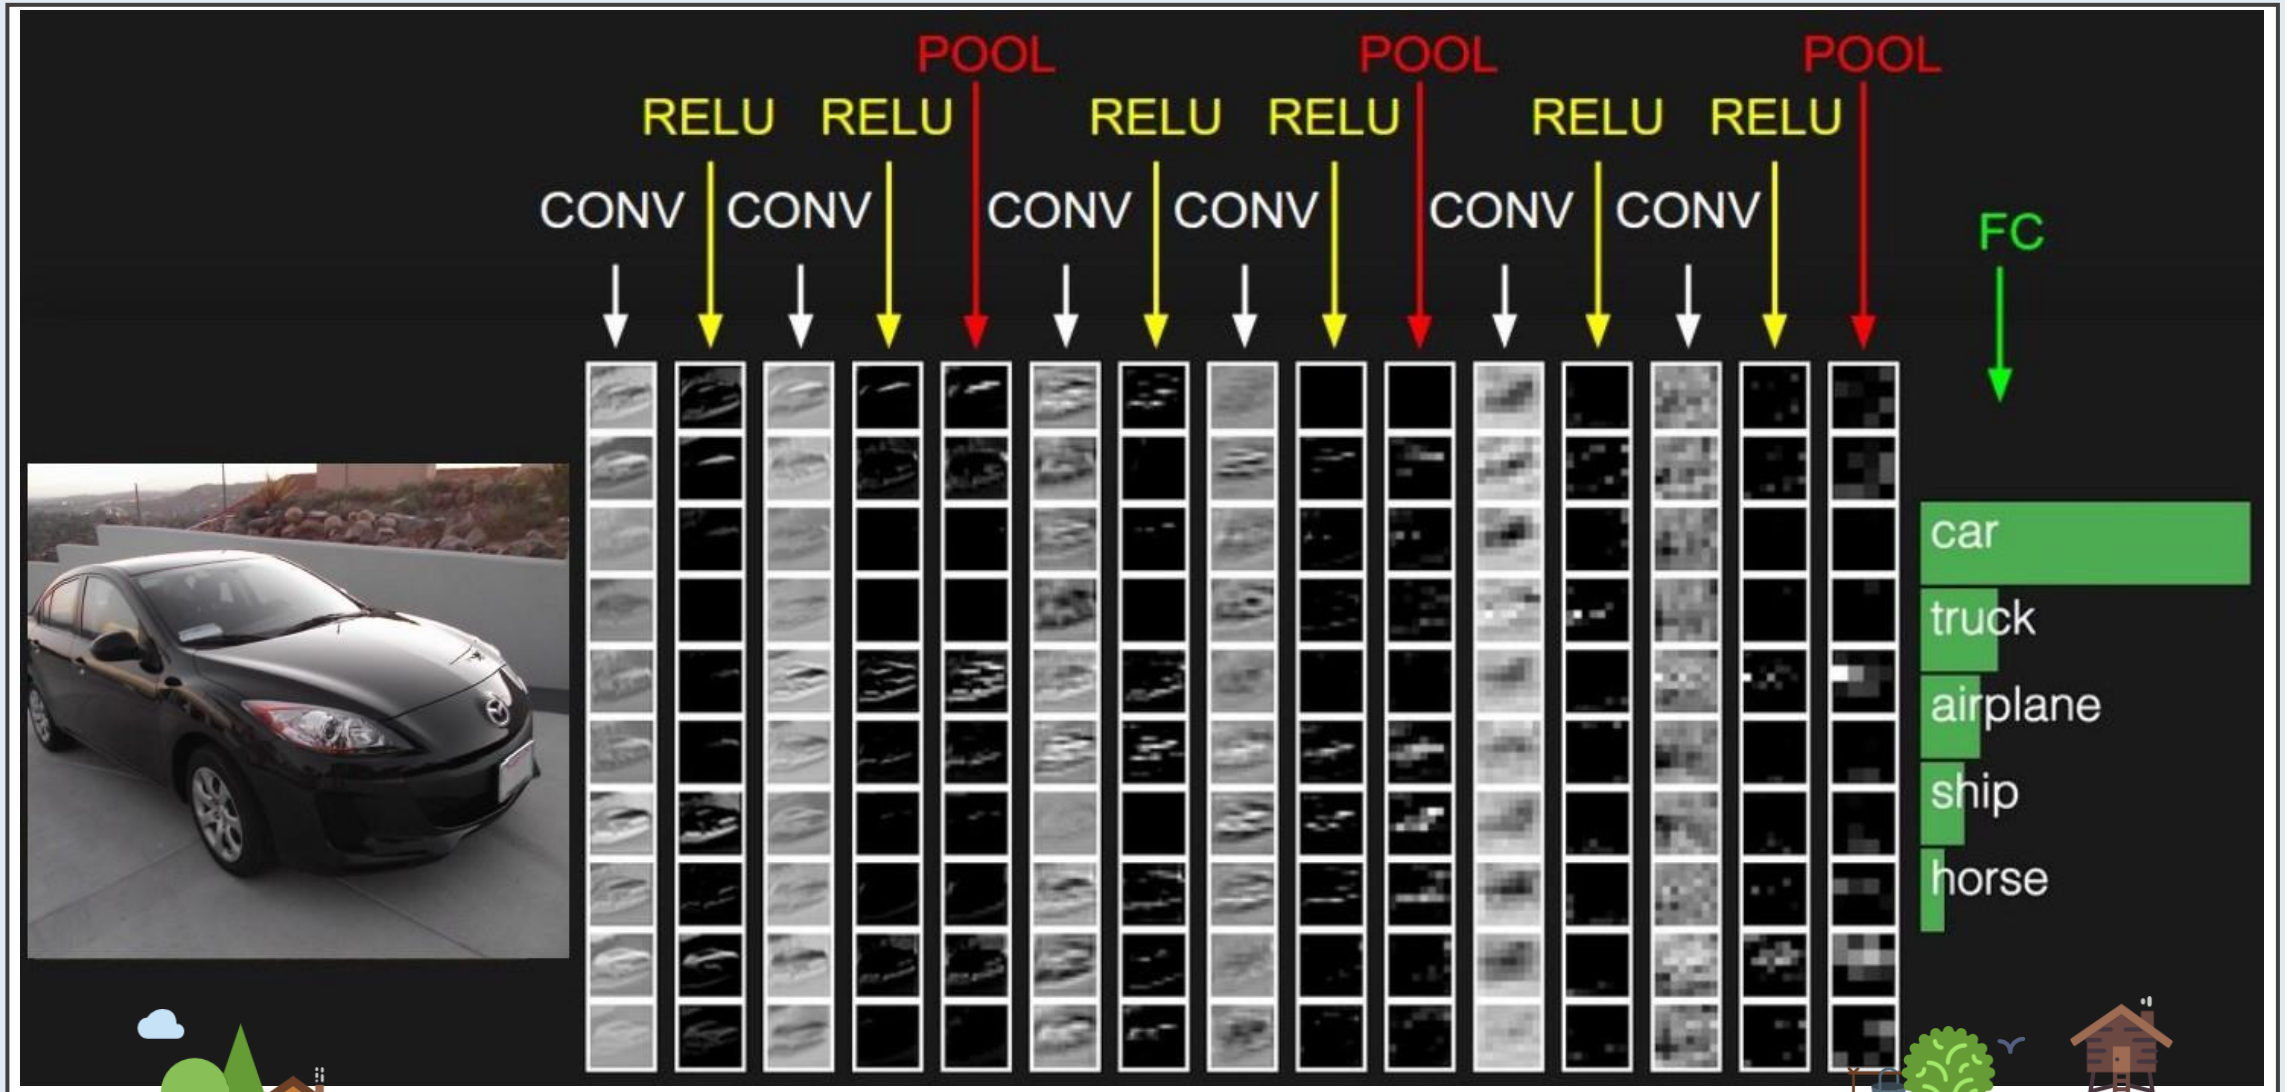

# Zero padding

**Preview:** ConvNet is a sequence of Convolution Layers, interspersed with activation functions

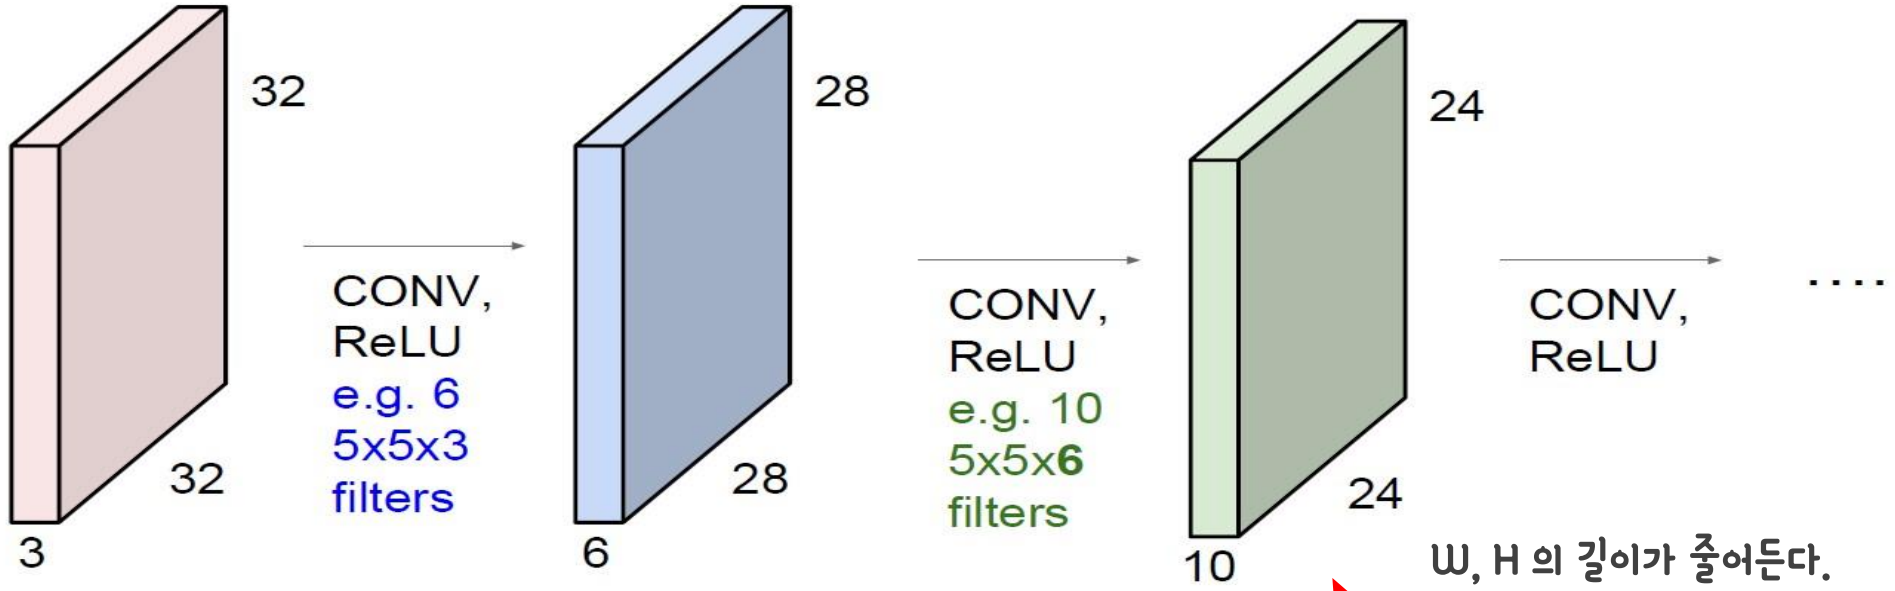


W, H 의 길이가 줄어든다.

해결책 : zero padding

# Zero padding

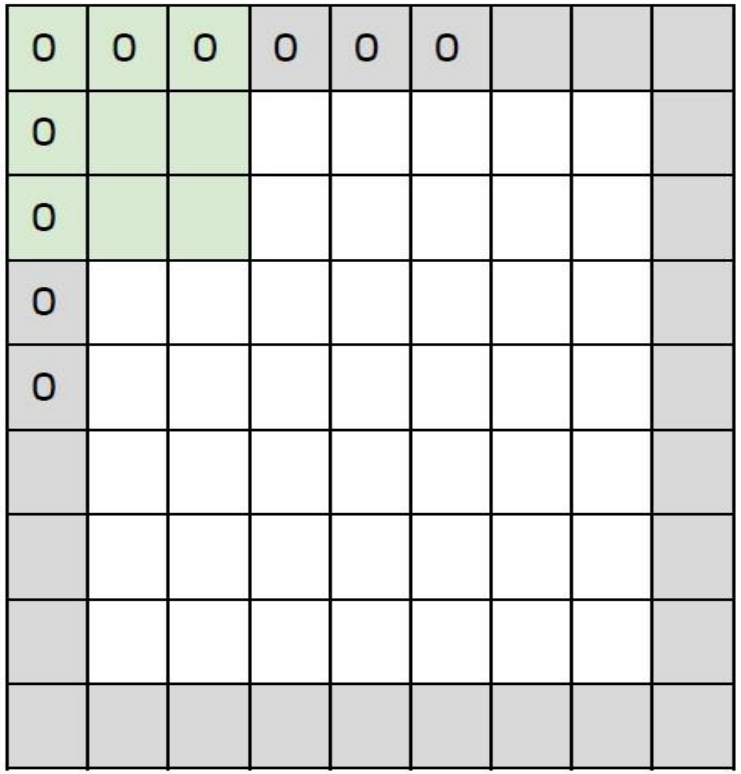


- Activation map size 조절
- 이 후 연산은 동일

# Example

Examples time:

Input volume: **32x32x3**
10 5x5 filters with stride 1, pad 2

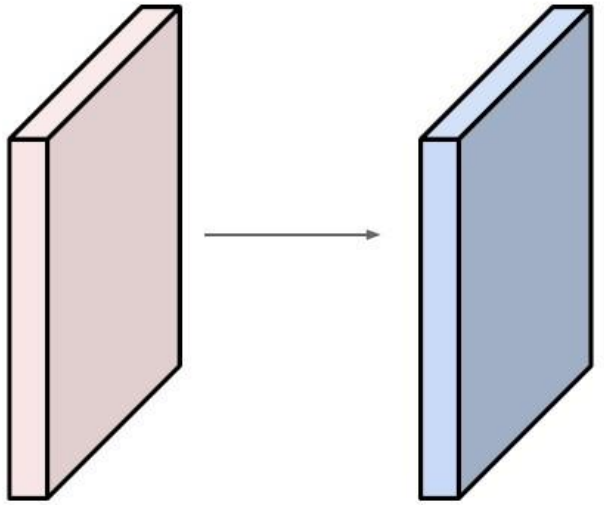

Output volume size: ?

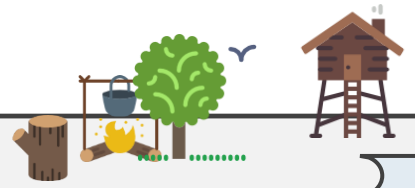

# Example

Examples time:

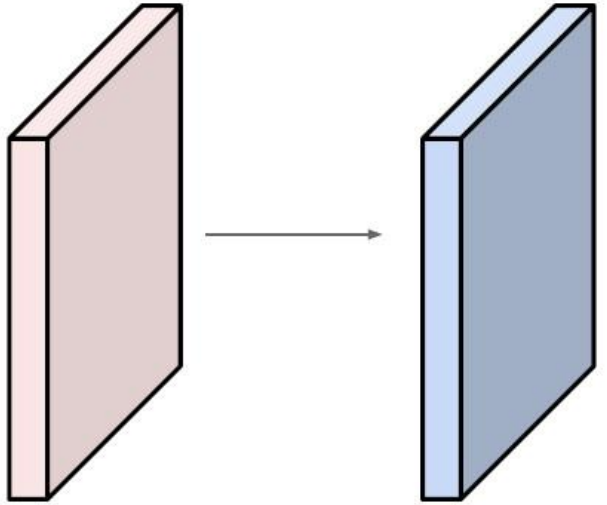

Input volume: **32x32x3**
10 5x5 filters with stride 1, pad 2

Output volume size:
(32+2*2-5)/1+1 = 32 spatially, so
**32x32x10**

Zero padding으로 인해 w, h 의size가 동일

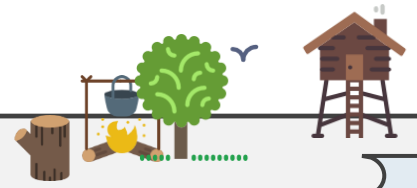

# Example

Examples time:

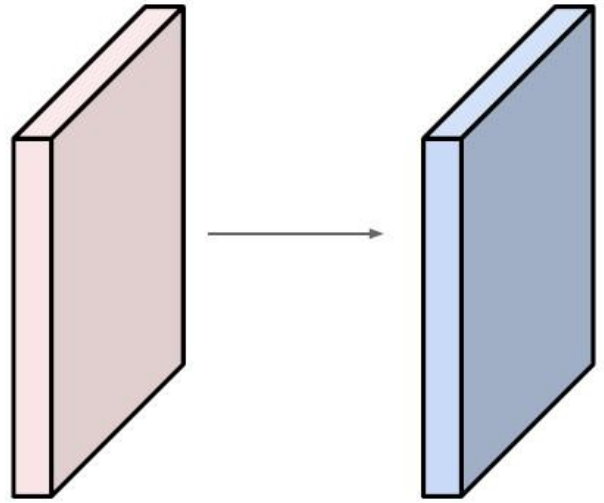

Input volume: **32x32x3**
10 5x5 filters with stride 1, pad 2

Number of parameters in this layer?

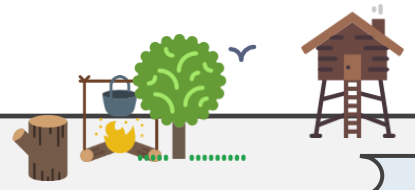

# Example

Examples time:

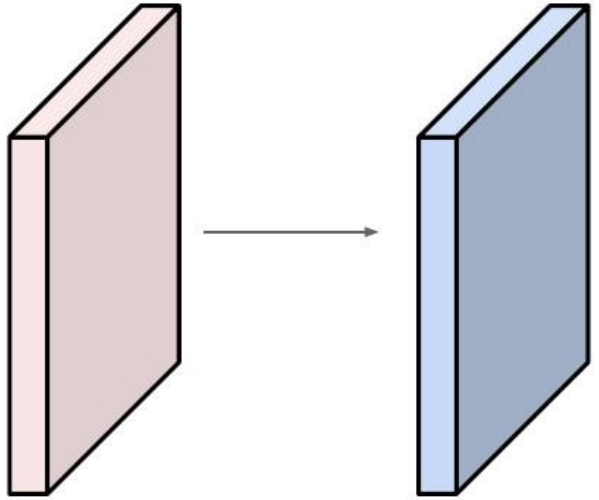

Input volume: **32x32x3**
10 5x5 filters with stride 1, pad 2

Number of parameters in this layer?
each filter has 5*5*3 + 1 = 76 params (+1 for bias)
=> 76*10 = **760**

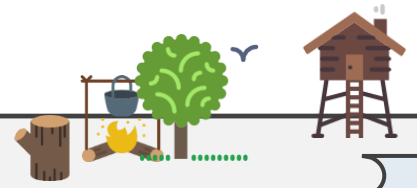

# Pooling layer

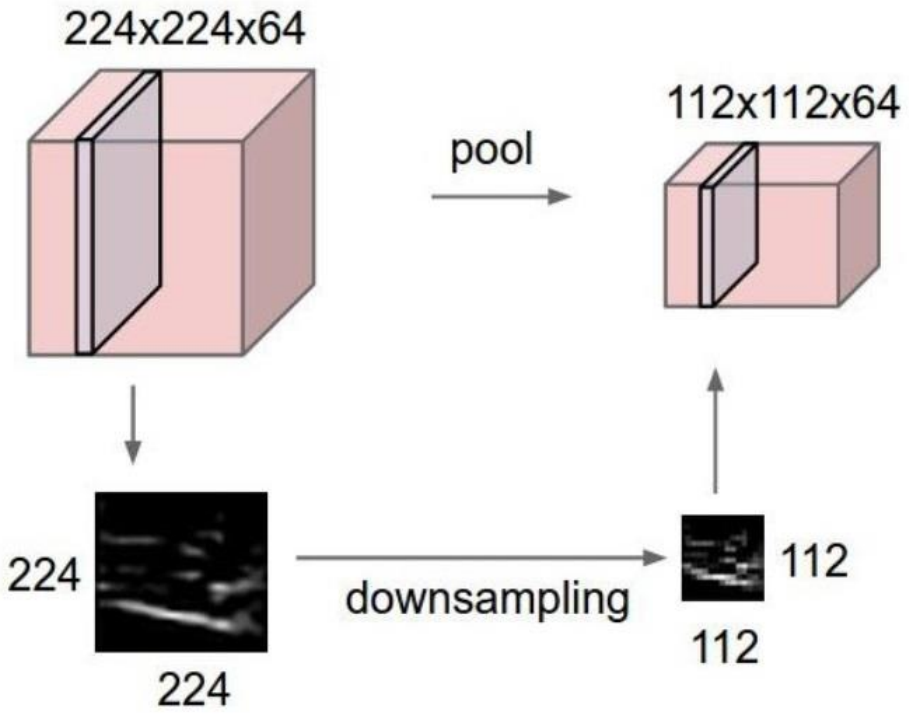

# Max pooling

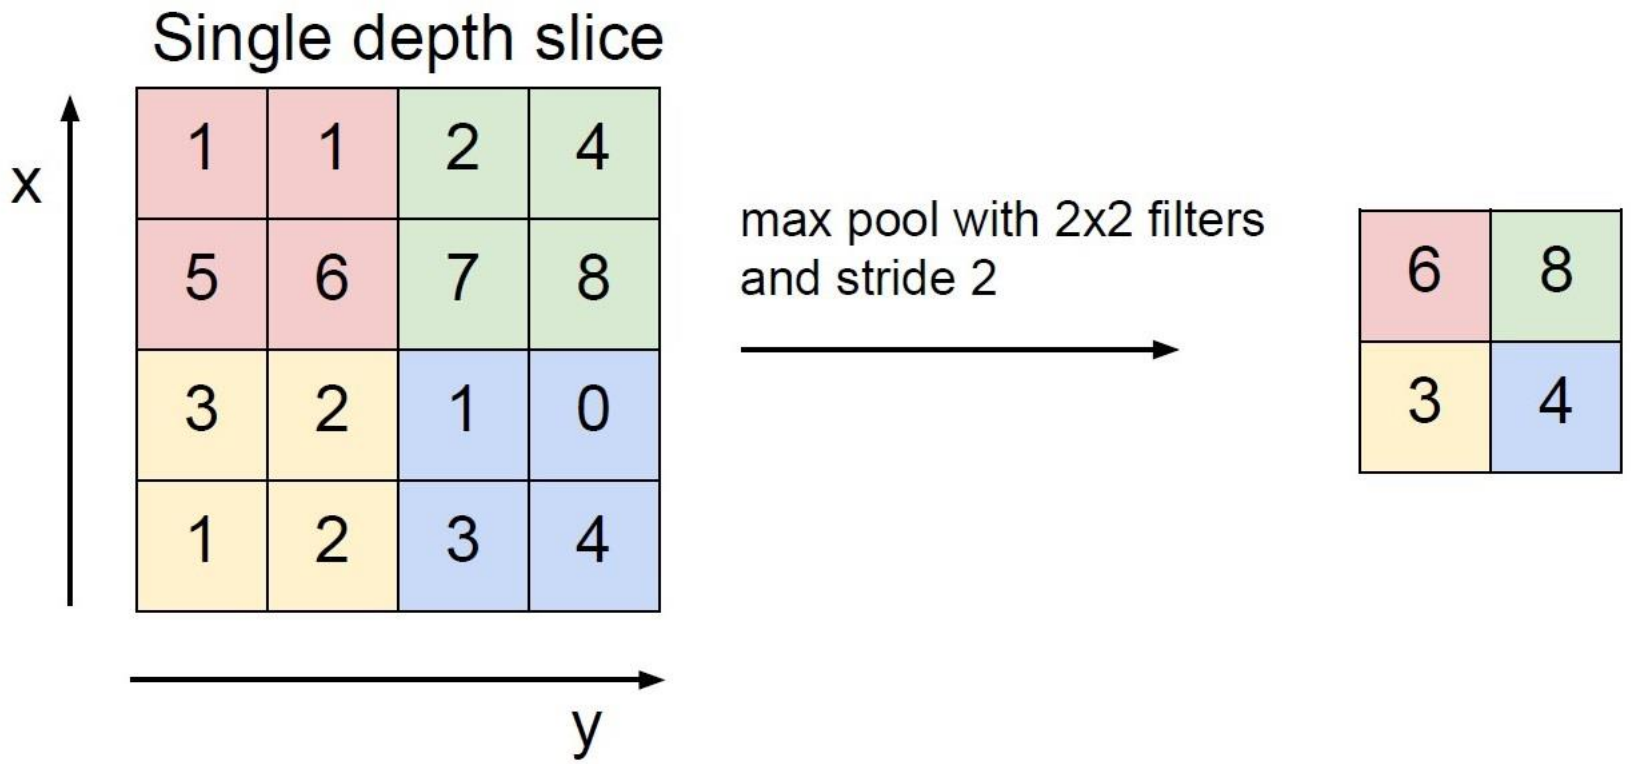
Single depth slice
x
y
1 1 2 4
5 6 7 8
3 2 1 0
1 2 3 4
max pool with 2x2 filters
and stride 2
6 8
3 4

# Average pooling

| | | | |
|---|---|---|---|
| 2 | 3 | 1 | 9 |
| 4 | 7 | 3 | 5 |
| 8 | 2 | 2 | 2 |
| 1 | 3 | 4 | 5 |

→

| | |
|---|---|
| 4 | 4.5 |
| 3.25 | 3.25 |

End

# 참고 자료

CS231n : http://cs231n.stanford.edu/syllabus.html